

## Chapter 2

# The Pandas essentials for data analysis

# Objectives (part 1)

## Applied

1. Create a DataFrame by reading data from a file or by using the DataFrame constructor.
2. Save a DataFrame to disk as a pickle file, and restore the DataFrame by reading the pickle file.
3. Examine the data in a DataFrame by displaying the data and its attributes.
4. Examine the data in a DataFrame by using the info(), nunique(), and describe() methods.
5. Access columns, rows, or a subset of columns and rows by using some combination of dot notation, brackets, the query() method, and the loc[] or iloc[] accessor.
6. Use Pandas methods to get statistics for the columns of a DataFrame.

## Objectives (part 2)

7. Use Python to do calculations on the data in the columns of a DataFrame.
8. Use the Pandas or Python replace() method to replace the data in a DataFrame or Series object.
9. Use Pandas methods to do the following DataFrame operations:
  - Sort the rows
  - Set and reset an index
  - Pivot the data
  - Melt the data
  - Group and aggregate the data
  - Plot the data based on the index that has been set

# Objectives (part 3)

## Knowledge

1. Distinguish between a Series and a DataFrame.
2. Describe these attributes of a DataFrame: values, index, columns, shape.

# The first five rows of the Child Mortality data in a DataFrame

	Year	Age Group	Death Rate
0	1900	1-4 Years	1983.8
1	1901	1-4 Years	1695.0
2	1902	1-4 Years	1655.7
3	1903	1-4 Years	1542.1
4	1904	1-4 Years	1591.5

# The components of a DataFrame

Component	Description
Column labels	The names at the tops of the columns.
Column data	The data in the columns. All of the data in a column typically has the same data type with one entry in each row.
Column data types	Each column has a defined data type. If all of the elements in a column don't have the same data type, the elements are stored with the object data type.
Index	Also known as a row label. If an index isn't defined, it is generated as a sequence of integers starting with zero.
Metadata	Attributes of the DataFrame that are generated by Pandas when the DataFrame is constructed or changed.

# Terms related to DataFrame objects

- DataFrame
- Column
- Row
- Series object
- Element
- Datapoint

# Three of the Pandas read() methods for importing data into a DataFrame

```
read_csv()  
read_excel()  
read_sql_query()
```

## How to import a CSV file from a website

```
import pandas as pd  
url = "https://data.cdc.gov/api/views/v6ab-adf5/" +  
      "rows.csv?accessType=DOWNLOAD"  
mortality_data = pd.read_csv(url)
```

# The DataFrame() constructor

`DataFrame(params)`

## Parameters

`data`

`columns`

`index`

## The data and columns arrays for a DataFrame

```
df_data=[[1900, '1-4 Years', 1983.8],[1901, '1-4 Years', 1695.0]]  
df_columns=['Year', 'Age Group', 'Death Rate']
```

## The code that creates the DataFrame

```
import pandas as pd  
mortality_df = pd.DataFrame(  
    data=df_data,  
    columns=df_columns)
```

## The DataFrame that's created

	Year	Age Group	Death Rate
0	1900	1-4 Years	1983.8
1	1901	1-4 Years	1695.0

# Three of the Pandas methods for saving a DataFrame to disk

- `to_pickle()`
- `to_csv()`
- `to_excel()`

## How to save a DataFrame as a pickle file

```
mortality_data.to_pickle('mortality_data.pkl')
```

## The Pandas `read_pickle()` method

```
read_pickle(filename)
```

### How to read a pickle file to restore a DataFrame

```
mortality_data = pd.read_pickle('mortality_data.pkl')
```

### The DataFrame that's restored

	Year	Age Group	Death Rate
0	1900	1-4 Years	1983.8
1	1901	1-4 Years	1695.0
2	1902	1-4 Years	1655.7
3	1903	1-4 Years	1542.1
4	1904	1-4 Years	1591.5

# Two methods for displaying the data in a DataFrame

`head(rows)`

`tail(rows)`

# JupyterLab automatically displays the contents of a named DataFrame

`mortality_data`

	Year	Age Group	Death Rate
0	1900	1-4 Years	1983.8
1	1901	1-4 Years	1695.0
2	1902	1-4 Years	1655.7
3	1903	1-4 Years	1542.1
4	1904	1-4 Years	1591.5
...	...	...	...
471	2014	15-19 Years	45.5
472	2015	15-19 Years	48.3
473	2016	15-19 Years	51.2
474	2017	15-19 Years	51.5
475	2018	15-19 Years	49.2
476 rows × 3 columns			

## How to use the head() and tail() methods

```
mortality_data.head()          # displays the first 5 rows  
mortality_data.tail(3)         # displays the last 3 rows
```

## How to display the data in 5 rows and all columns

```
with pd.option_context(  
    'display.max_rows', 5,  
    'display.max_columns', None):  
    display(mortality_data)
```

# Some of the attributes of a DataFrame object

Attribute	Description
<code>values</code>	The values of the DataFrame in an array format
<code>index</code>	The row index
<code>columns</code>	The column names
<code>size</code>	The total number of elements
<code>shape</code>	The number of rows and columns

# The values attribute

```
mortality_data.values  
=====array([[1900, '1-4 Years', 1983.8],  
           [1901, '1-4 Years', 1695.0],  
           [1902, '1-4 Years', 1655.7],  
           ...,  
           [2016, '15-19 Years', 51.2],  
           [2017, '15-19 Years', 51.5],  
           [2018, '15-19 Years', 49.2]], dtype=object)
```

## The other four attributes

```
print("Index:  ", mortality_data.index)
print("Columns:", mortality_data.columns)
print("Size:   ", mortality_data.size)
print("Shape:  ", mortality_data.shape)
=====
Index:  RangeIndex(start=0, stop=476, step=1)
Columns: Index(['Year', 'Age Group', 'Death Rate'], dtype='object')
Size:   1428
Shape:  (476, 3)
```

# How to use the columns attribute to replace spaces with nothing

```
mortality_data.columns =  
    mortality_data.columns.str.replace(' ', '')
```

## The new column names

```
Index(['Year', 'AgeGroup', 'DeathRate'], dtype='object')
```

# The `info()`, `nunique()`, and `describe()` methods

Method	Description
<code>info(params)</code>	Returns information about the DataFrame and its columns.
<code>nunique()</code>	Returns the number of unique data items in each column.
<code>describe()</code>	Returns statistical information for each numeric column.

## How to use the info() method

```
mortality_data.info()  
=====
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 476 entries, 0 to 475  
Data columns (total 3 columns):  
Year 476 non-null int64  
AgeGroup 476 non-null object  
DeathRate 476 non-null float64  
dtypes: float64(1), int64(1), object(1)  
memory usage: 11.3+ KB

## The memory\_usage parameter ensures a more accurate usage result

```
mortality_data.info(memory_usage='deep')  
=====
```

<class 'pandas.core.frame.DataFrame'>  
...  
memory usage: 38.7 KB

## How to use the nunique() method

```
mortality_data.nunique()  
=====  
Year          119  
AgeGroup      4  
DeathRate     430  
dtype: int64
```

# How to use the describe() method

**With the T property, the statistics are displayed in the columns**

```
mortality_data.describe().T
```

	count	mean	std	min	25%	50%	75%	max
Year	476.0	1959.00000	34.387268	1900.0	1929.000	1959.0	1989.000	2018.0
DeathRate	476.0	192.92416	293.224216	11.4	40.575	89.5	222.575	1983.8

**Without the T property, the statistics are displayed in the rows**

```
mortality_data.describe()
```

# The syntax for accessing columns

To access	With brackets	With dot notation
One column	<code>df[column_name]</code> <code>df[[column_name]]</code>	<code>df.column_name</code>
Two or more columns	<code>df[[col1,col2,...]]</code>	

# Two ways to access one column

## With dot notation

```
mortality_data.DeathRate.head(2)
```

```
=====
```

```
0    1983.8
```

```
1    1695.0
```

```
Name: DeathRate, dtype: float64
```

## With brackets

```
mortality_data['DeathRate'].head(2)
```

```
=====
```

```
0    1983.8
```

```
1    1695.0
```

```
Name: DeathRate, dtype: float64
```

## How to use a list to access two or more columns

```
mortality_data[['Year', 'DeathRate']].head(2)
```

	Year	DeathRate
0	1900	1983.8
1	1901	1695.0

# The query() method

`query(condition)`

## How to use the query() method to access rows (part 1)

Based on a single column

```
mortality_data.query('Year == 1900')
```

	Year	AgeGroup	DeathRate
0	1900	1-4 Years	1983.8
119	1900	5-9 Years	466.1
238	1900	10-14 Years	298.3
357	1900	15-19 Years	484.8

# How to use the query() method to access rows (part 2)

Based on multiple columns with the *and* operator

```
mortality_data.query('Year == 2000 and AgeGroup != "1-4 Years"')
```

	Year	AgeGroup	DeathRate
219	2000	5-9 Years	15.8
338	2000	10-14 Years	20.3
457	2000	15-19 Years	67.1

# How to use the query() method to access rows (part 3)

Based on multiple columns with the `or` operator

```
mortality_data.query('Year == 1900 or Year == 2000').head()
```

	Year	AgeGroup	DeathRate
0	1900	1-4 Years	1983.8
100	2000	1-4 Years	32.4
119	1900	5-9 Years	466.1
219	2000	5-9 Years	15.8
238	1900	10-14 Years	298.3

With backticks for column names that contain spaces

```
mortality_data.query('Year == 2000 and `Age Group` != "1-4 Years"')
```

## How to access one column from a subset of rows using dot notation

```
mortality_data.query('Year == 1900').DeathRate  
=====  
0      1983.8  
119    466.1  
238    298.3  
357    484.8  
Name: DeathRate, dtype: float64
```

## How to access one column from a subset of rows using brackets

```
mortality_data.query('Year == 1900')['DeathRate']  
=====  
0      1983.8  
119    466.1  
238    298.3  
357    484.8  
Name: DeathRate, dtype: float64
```

# How to access one column from a subset of rows using a list

```
mortality_data.query('Year == 1900')[['DeathRate']]
```

DeathRate	
0	1983.8
119	466.1
238	298.3
357	484.8

# How to access two or more columns from a subset of rows using a list

```
mortality_data.query('Year == 1900')[['AgeGroup', 'DeathRate']]
```

	AgeGroup	DeathRate
0	1-4 Years	1983.8
119	5-9 Years	466.1
238	10-14 Years	298.3
357	15-19 Years	484.8

# The `loc[ ]` and `iloc[ ]` accessors for accessing rows and columns

Accessor	Accesses rows and columns
<code>loc[rows,columns]</code>	by their labels.
<code>iloc[rows,columns]</code>	by their positions.

# How to access rows with the loc[ ] accessor

**With a list to access the rows with labels 0, 5, and 10**

```
mortality_data.loc[[0,5,10]]
```

**With a slice to access the rows with labels 4 through 6**

```
mortality_data.loc[4:6]
```

**With a slice to access every 5th row from 0 through 20**

```
mortality_data.loc[0:20:5]
```

**With a conditional expression to access the rows  
for the year 1917**

```
mortality_data.loc[mortality_data.Year == 1917]
```

# How to access rows and columns with the loc[ ] accessor

## With lists of row and column labels

```
mortality_data.loc[[0,5,10],['AgeGroup','DeathRate']]
```

## With slices of row and column labels

```
mortality_data.loc[4:6,'AgeGroup':'DeathRate']
```

# How to access rows and columns with the `iloc[ ]` accessor

## With lists of row and column positions

```
mortality_data.iloc[[4,5,6],[1,2]]
```

## With slices of row and column positions

```
mortality_data.iloc[4:7,1:3]
```

## With a negative row position in a slice to access the last 10 rows

```
mortality_data.iloc[-10:]
```

## The `sort_values()` method

```
sort_values(columns, ascending)
```

### How to sort by one column in descending order

```
mortality_data.sort_values('DeathRate',  
                           ascending=False).head(3)
```

	Year	AgeGroup	DeathRate
0	1900	1-4 Years	1983.8
1	1901	1-4 Years	1695.0
2	1902	1-4 Years	1655.7

# How to sort by multiple columns in ascending order

```
mortality_data.sort_values(['Year', 'DeathRate']).head(3)
```

	Year	AgeGroup	DeathRate
238	1900	10-14 Years	298.3
119	1900	5-9 Years	466.1
357	1900	15-19 Years	484.8

## How to sort by multiple columns in mixed orders

```
mortality_data.sort_values(['Year', 'DeathRate'],  
                           ascending=[True, False]).head()
```

	Year	AgeGroup	DeathRate
0	1900	1-4 Years	1983.8
357	1900	15-19 Years	484.8
119	1900	5-9 Years	466.1
238	1900	10-14 Years	298.3
1	1901	1-4 Years	1695.0

# Some of the Pandas methods for both Series and DataFrame objects

`count()`

`mean()`

`median()`

`min()`

`max()`

`std()`

`sum()`

`cumsum()`

`quantile(q)`

## How to apply a method to one column

```
mortality_data.DeathRate.mean()  
=====  
192.92415966386568
```

## How to apply a method to two columns

```
mortality_data[['AgeGroup', 'DeathRate']].max()  
=====  
AgeGroup      5-9 Years  
DeathRate     1983.8  
dtype: object
```

## How to apply a method to all columns

```
mortality_data.count()  
=====  
Year        476  
AgeGroup    476  
DeathRate   476  
dtype: int64
```

# How to apply the quantile() method to two different quantiles

```
mortality_data.quantile([.1,.9])
```

	Year	DeathRate
0.1	1911.5	21.50
0.9	2006.5	430.85

# The Python operators for column arithmetic

Operator	Name
+	Addition
-	Subtraction
*	Multiplication
/	Division
//	Integer division
%	Modulo/Remainder
**	Exponentiation

## How to add a column to a DataFrame

```
mortality_data['MeanCentered'] = \
    mortality_data.DeathRate - mortality_data.DeathRate.mean()
```

	Year	AgeGroup	DeathRate	MeanCentered
0	1900	1-4 Years	1983.8	1790.87584
1	1901	1-4 Years	1695.0	1502.07584
2	1902	1-4 Years	1655.7	1462.77584
3	1903	1-4 Years	1542.1	1349.17584

## How to modify the data in an existing column

```
mortality_data[ 'DeathRate' ] = \  
    mortality_data.DeathRate / 100000
```

	Year	AgeGroup	DeathRate	MeanCentered
0	1900	01-04 Years	0.019838	1790.87584
1	1901	01-04 Years	0.016950	1502.07584
2	1902	01-04 Years	0.016557	1462.77584
3	1903	01-04 Years	0.015421	1349.17584

## The Pandas replace() method

```
replace(to_replace,value, inplace)
```

## The Python replace() method

```
replace(old,new)
```

# How to modify the string data in a column

## With the Pandas replace() method with three parameters

```
mortality_data.AgeGroup.replace(  
    to_replace = ['1-4 Years', '5-9 Years'],  
    value = ['01-04 Years', '05-09 Years'],  
    inplace = True)
```

## With the Pandas replace() method and a dictionary of old and new values

```
mortality_data.AgeGroup.replace(  
    {'1-4 Years': '01-04 Years', '5-9 Years': '05-09 Years'},  
    inplace = True)
```

## With the Python replace() method

```
mortality_data[ 'AgeGroup' ] =  
    mortality_data.AgeGroup.str.replace('1-4 Years', '01-04 Years')  
mortality_data[ 'AgeGroup' ] =  
    mortality_data.AgeGroup.str.replace('5-9 Years', '05-09 Years')
```

## The result for all three of the examples

	Year	AgeGroup	DeathRate	MeanCentered
0	1900	01-04 Years	0.019838	1790.87584
1	1901	01-04 Years	0.016950	1502.07584
2	1902	01-04 Years	0.016557	1462.77584
3	1903	01-04 Years	0.015421	1349.17584

## The `set_index()` method

```
set_index(columns, verify_integrity)
```

### How to set a one-column index

```
mortality_data = mortality_data.set_index('Year')  
mortality_data.head(2)
```

	AgeGroup	DeathRate	MeanCentered
Year			
1900	01-04 Years	0.019838	1790.87584
1901	01-04 Years	0.016950	1502.07584

## How to verify the integrity of a one-column index

```
mortality_data = mortality_data.set_index('Year', verify_integrity=True)
mortality_data.head(2)
=====
ValueError: Index has duplicate keys: Int64Index([1900, 1901, ....
```

## How to set a two-column index

```
mortality_data.set_index(['Year', 'AgeGroup'],  
    verify_integrity=True)  
mortality_data.head(2)
```

		DeathRate	MeanCentered
Year	AgeGroup		
1900	01-04 Years	0.019838	1790.87584
1901	01-04 Years	0.016950	1502.07584

# The `reset_index()` method

```
reset_index(inplace)
```

## How to reset an index

```
mortality_data.reset_index(inplace=True)  
mortality_data.head(2)
```

	Year	AgeGroup	DeathRate	MeanCentered
0	1900	01-04 Years	0.019838	1790.87584
1	1901	01-04 Years	0.016950	1502.07584

# The pivot() method of a DataFrame

```
pivot(index,columns,values)
```

## How to pivot the AgeGroup column for the death rate

```
mortality_wide = mortality_data.pivot(  
    index='Year', columns='AgeGroup', values='DeathRate')  
mortality_wide.head(3)
```

AgeGroup	01-04 Years	05-09 Years	10-14 Years	15-19 Years
Year				
1900	0.019838	0.004661	0.002983	0.004848
1901	0.016950	0.004276	0.002736	0.004544
1902	0.016557	0.004033	0.002525	0.004215

# How to pivot the AgeGroup column for all other data

```
mortality_wide = mortality_data.pivot(  
    index='Year', columns='AgeGroup')  
mortality_wide.head(3)
```

AgeGroup	DeathRate					MeanCentered		
	01-04 Years	05-09 Years	10-14 Years	15-19 Years	01-04 Years	05-09 Years	10-14 Years	15-19 Years
Year								
1900	0.019838	0.004661	0.002983	0.004848	1790.87584	273.17584	105.37584	291.87584
1901	0.016950	0.004276	0.002736	0.004544	1502.07584	234.67584	80.67584	261.47584
1902	0.016557	0.004033	0.002525	0.004215	1462.77584	210.37584	59.57584	228.57584

# The melt() method of a DataFrame

Method	Description
<code>melt(params)</code>	Melts the data in two or more columns into a single column.

## Parameters of the melt() method

Parameters	Description
<code>id_vars</code>	The column or columns that identify each row.
<code>value_vars</code>	The columns to melt. If none are specified, all other columns will be melted.
<code>var_name</code>	The name of the column that will contain the melted column names, or “variable” by default.
<code>value_name</code>	The name of the column that will contain the values, or “value” by default.

# The starting DataFrame in wide form

	Year	01-04 Years	05-09 Years	10-14 Years	15-19 Years
0	1900	0.019838	0.004661	0.002983	0.004848
1	1901	0.016950	0.004276	0.002736	0.004544
2	1902	0.016557	0.004033	0.002525	0.004215
3	1903	0.015421	0.004147	0.002682	0.004341

# How to melt the data for just the 01-04 and 05-09 columns

```
mortality_long = mortality_wide.melt(  
    id_vars = 'Year',  
    value_vars=[ '01-04 Years', '05-09 Years' ],  
    var_name = 'AgeGroup',  
    value_name= 'DeathRate')
```

	Year	AgeGroup	DeathRate
0	1900	01-04 Years	0.019838
1	1901	01-04 Years	0.016950
...	...	...	...
236	2017	05-09 Years	0.000116
237	2018	05-09 Years	0.000115

# The groupby() method

groupby(columns)

## How to group the data on the AgeGroup column

```
mortality_data.groupby('AgeGroup').mean().head(4)
```

AgeGroup	Year	DeathRate	MeanCentered
01-04 Years	1959	0.003832	190.301891
05-09 Years	1959	0.001173	-75.598109
10-14 Years	1959	0.000938	-99.154412
15-19 Years	1959	0.001774	-15.549370

## How to group the data on the Year column

```
mortality_data.groupby('Year').median().head(4)
```

	DeathRate	MeanCentered
Year		
1900	0.004755	282.52584
1901	0.004410	248.07584
1902	0.004124	219.47584
1903	0.004244	231.47584

## How to group the data on multiple columns

```
mortality_data.groupby(['Year', 'AgeGroup']).count().head()
```

Year	AgeGroup	DeathRate	MeanCentered
1900	01-04 Years	1	1
	05-09 Years	1	1
	10-14 Years	1	1
	15-19 Years	1	1
1901	01-04 Years	1	1

# The agg() method

`agg(methods)`

## How to aggregate the data for all columns in each age group

```
mortality_data.groupby('AgeGroup').agg(['mean','median'])
```

AgeGroup	Year		DeathRate		MeanCentered	
	mean	median	mean	median	mean	median
01-04 Years	1959	1959	0.003832	0.001091	190.301891	-83.82416
05-09 Years	1959	1959	0.001173	0.000484	-75.598109	-144.52416
10-14 Years	1959	1959	0.000938	0.000446	-99.154412	-148.32416
15-19 Years	1959	1959	0.001774	0.001069	-15.549370	-86.02416

# How to aggregate the data for just the death rate in each age group

```
mortality_data.groupby('AgeGroup')[ 'DeathRate' ] \  
.agg(['mean','median','std','nunique'])
```

	mean	median	std	nunique
AgeGroup				
01-04 Years	0.003832	0.001091	0.005005	117
05-09 Years	0.001173	0.000484	0.001275	115
10-14 Years	0.000938	0.000446	0.000884	115
15-19 Years	0.001774	0.001069	0.001384	117

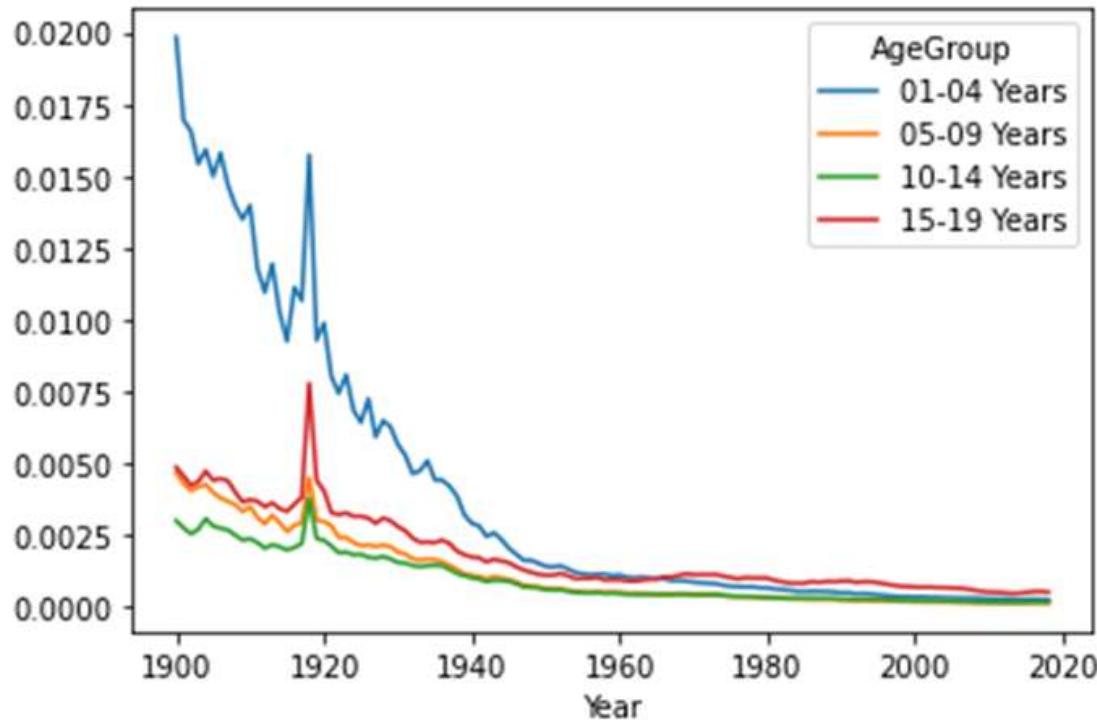
# How to aggregate the data for just the death rate in each year group

```
mortality_data.groupby('Year')[['DeathRate']] \\\n    .agg(['mean','median','std','min','max','var','nunique'])
```

	mean	median	std	min	max	var	nunique
Year							
1900	0.008082	0.004755	0.007882	0.002983	0.019838	6.212178e-05	4
1901	0.007127	0.004410	0.006597	0.002736	0.016950	4.352410e-05	4
1902	0.006832	0.004124	0.006527	0.002525	0.016557	4.260299e-05	4

## How to chain the pivot() and plot() methods

```
mortality_data.pivot(index='Year',  
columns='AgeGroup')['DeathRate'].plot()
```



# How to chain the groupby(), agg(), and plot() methods

```
mortality_data.groupby('AgeGroup')['DeathRate'] \  
    .agg(['mean','median','std']).plot.barh()
```

